# KAVA: LATENT REASONING VIA COMPRESSED KV-CACHE DISTILLATION

Anna Kuzina\* Qualcomm AI Research† akuzina@qti.qualcomm.com

Paul N. Whatmough Qualcomm AI Research pwhatmou@qti.qualcomm.com Maciej Pioro\*†
IDEAS NCBR / IPPT PAN
maciej.pioro@gmail.com

Babak Ehteshami Bejnordi Qualcomm AI Research behtesha@qti.qualcomm.com

#### **ABSTRACT**

Large Language Models (LLMs) excel at multi-step reasoning problems with explicit chain-of-thought (CoT), but verbose traces incur significant computational costs and memory overhead, and often carry redundant, stylistic artifacts. Latent reasoning has emerged as an efficient alternative that internalizes the thought process, but it suffers from a critical lack of supervision, limiting its effectiveness on complex, natural-language reasoning traces. In this work we propose KAVA, the first framework that bridges this gap by distilling knowledge directly from a compressed KV-cache of the teacher into a latent-reasoning student via selfdistillation, leveraging the representational flexibility of continuous latent tokens to align stepwise KV trajectories. We show that the abstract, unstructured knowledge within compressed KV-cache, which lacks direct token correspondence, can serve as a rich supervisory signal for a latent reasoning student. Empirically, the approach consistently outperforms strong latent baselines, exhibits markedly smaller degradation from equation-only to natural-language traces, and scales to larger backbones while preserving efficiency. These results establish compressed KV-cache distillation as a scalable supervision signal for latent reasoning, combining the accuracy of CoT-trained teachers with the efficiency and deployability of latent inference.

### 1 Introduction

Recent advancements in Large Language Models (LLMs) have demonstrated remarkable capabilities in solving complex problems across domains such as mathematics (Zhang et al., 2025), science (Phan et al., 2025), and code generation (Hui et al., 2024). A key driver of this progress is "chain-of-thought" (CoT) training that elicits intermediate steps before the final answer, improving accuracy on long-horizon inference problems (DeepSeek-AI et al., 2025). Yet, explicit CoT often incurs substantial inference cost due to long, verbose traces and the associated key-value (KV) cache growth, making deployment on memory- and compute-constrained devices difficult. Furthermore, CoT traces, especially those distilled from larger models, can inherit and amplify biases or contain plausible-sounding but fallacious logic, limiting their reliability.

Recent studies show that the KV-caches underlying CoT are highly redundant and can be aggressively compressed with little to no loss in accuracy (Cai et al., 2025; Park et al., 2025), indicating that much of CoT's signal resides in compressible structure rather than indispensable text. This observation suggests an alternative supervisory path: if the essential dynamics of reasoning live in the cache, perhaps models can be trained to internalize those dynamics without verbose traces at inference time. However, this compressed KV-cache presents a significant challenge for knowledge

<sup>\*</sup>Equal contribution

<sup>&</sup>lt;sup>†</sup>Qualcomm AI Research is an initiative of Qualcomm Technologies, Inc.

<sup>&</sup>lt;sup>‡</sup>Work done during internship at Qualcomm AI Research.

distillation. As pruning decisions are often made independently per layer and attention head, the resulting compressed KV vectors lose their direct correspondence to specific input tokens, rendering conventional distillation schemes that match token activations or layer-wise hidden states ill-posed and non-trivial.

Latent reasoning is a nascent but promising direction in which reasoning occurs within the model's continuous latent space rather than being explicitly externalized (Hao et al., 2024; Su et al., 2025). Latent approaches promise efficiency by reducing token generation and KV-cache footprint, potentially closing the gap between strong reasoning performance and deployability in constrained settings. However, current latent reasoning methods struggle with the absence of direct supervision for internal thoughts, and successes are often reported in restricted setups; performance can degrade when training data contain long, natural-language-style traces that better reflect real-world reasoning workloads. In particular, compared to shorter, template-like traces, models trained on longer, natural-language reasoning sequences exhibit more fragile internal readouts and weaker generalization (Shen et al., 2025; Wu et al., 2025).

In this work, we bridge these gaps by introducing a novel framework that, for the first time, successfully distills the rich, abstract knowledge from a compressed teacher KV-cache into a latent reasoning student. We posit that the continuous, high-dimensional nature of latent thoughts provides a unique representational power that can absorb abstract cache structure that cannot be aligned at the token level. Concretely, our method is composed of three components: (i) the backbone that alternates between a teacher mode that consumes a full CoT to build per-layer, per-head KV-caches and a student mode that generates continuous latent thoughts; (ii) a redundancy- and importance-aware eviction module that compresses the teacher cache to the latent budget; (iii) and a KV-matching loss aligns the student's per-step latent K and V to the compressed target throughout the stack. This yields a strong, stepwise internal supervision signal that teaches the student to "think like" a compact cache of its own explicit reasoning while preserving the inference-time efficiency of latent reasoning. By supervising the latent trajectory directly in KV space, the approach bridges the gap between template-like latent traces and natural-language reasoning, yielding strong gains on natural-language datasets and scaling smoothly to larger backbones while retaining the efficiency benefits of latent inference. Our primary contributions are:

- We are the first to demonstrate that knowledge can be successfully distilled from a compressed KV-cache via self-distillation, despite the cache's head-wise, layer-wise eviction that destroys token correspondence.
- We show that by using the compressed KV-cache as a rich, step-by-step supervision signal, we can effectively train latent reasoners to learn directly from natural language traces where prior methods struggle to extract meaningful improvements.
- Through empirical evaluations, we show that our approach consistently outperforms strong latent baselines on natural language settings, exhibits smaller degradation when moving from equation-only to natural-language traces, and scales to larger backbones.

#### 2 BACKGROUND AND RELATED WORKS

Latent Reasoning. Traditional reasoning LLMs often rely on generating explicit intermediate steps in language to solve complex reasoning tasks. Recent work shifts reasoning from discrete text tokens to latent continuous tokens, where models perform iterative computation internally without generating external text (Chen et al., 2025; Zhu et al., 2025). Early work validated the benefit of extra computation through unstructured means, such as learnable pause tokens (Goyal et al., 2024) or even semantically meaningless filler tokens (Pfau et al., 2024), which improved performance on reasoning tasks by simply extending the model's processing time implicitly. Building on this implicit-compute view, iCoT moves from explicit to implicit CoT via distillation (Deng et al., 2023) and curriculum (Deng et al., 2024), progressively removing CoT while aligning internal states around answer prediction. This allows the model to internalize reasoning without generating text rationales at inference. Coconut (Hao et al., 2024) introduces "continuous thought" by feeding the last hidden state directly as the next input embedding, showing breadth-first search-like parallel exploration and fewer thinking tokens versus CoT on logical reasoning tasks. Follow-ups refine supervision and training dynamics: CODI (Shen et al., 2025) compresses CoT into continuous representations via self-distillation that supervises endpoints rather than full trajectories, while PCCoT (Wu et al.,

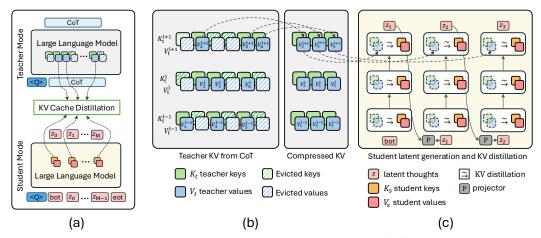


Figure 1: We propose KAVA, a latent reasoning model with KV-cache distillation loss. (a) Overview of our proposed compressed KV-cache distilled latent reasoning framework. (b) Teacher builds full KV-cache from a ground-truth CoT trace; a compression module produces a compressed cache to match the length of the latent trace; (c) a latent-reasoning student generates continuous thoughts  $z_t$  and is trained to match compressed teacher KV at each layer/step via KV distillation.

2025) parallelizes latent updates with Jacobi-style iterations to refine multiple continuous thoughts in tandem. In contrast to endpoint- or token-level supervision, our proposed approach distills a CoT teacher's compressed KV-cache into the student's latent trajectory, providing stepwise internal guidance that bridges the supervision gap in continuous-token reasoning without relying on explicit CoT text.

Complementary directions emphasize soft or hybrid traces: SoftCoT (Xu et al., 2025) injects soft thought tokens projected into the backbone's representation space to improve reasoning without altering hard-token generation, and Token Assorted (Su et al., 2025) mixes latent discrete tokens produced by a VQ-VAE with text tokens to shorten traces while maintaining accuracy. Our method is orthogonal, addressing the core challenge in latent reasoning, the absence of a direct supervision signal for these internal thoughts.

KV-cache Compression. KV-cache compression for reasoning focuses on trimming long, redundant thinking while preserving accuracy and throughput. R-KV (Cai et al., 2025) compresses onthe-fly by jointly scoring importance and redundancy to retain near-full performance with roughly 10–30% of the KV-cache on math reasoning, while KeyDiff (Park et al., 2025) offers a key-similarity-based eviction rule that preserves salient semantics under tight budgets. Other strategies such HeadKV (Fu et al., 2025), PyramidKV (Cai et al., 2024), LESS (Dong et al., 2024), and Eigen Attention (Saxena et al., 2024), provide complementary reductions via head selection, hierarchical/pyramidal retention, importance-aware mixed-precision, and low-rank attention, yielding robust long-context and reasoning behavior. KV-Distill (Chari et al., 2025) instead learns a lightweight adaptor that compresses long-context KV-caches and trains a compressed-cache student to match a full-cache teacher via output-level KL alignment. In contrast, our proposed approach treats the teacher's compressed KV-cache as supervision targets and distills them directly into the student's latent reasoning steps, aligning internal KV trajectories across the thinking process and directly addressing the lack of supervision for continuous thoughts.

#### 3 KAVA: KV-CACHE DISTILLATION FOR LATENT REASONING

#### 3.1 Overview

We will split the common chat template into three parts named question Q, reasoning trace C and answer A, with  $N_Q$ ,  $N_C$  and  $N_A$  token correspondingly. Consider an autoregressive generative model (LLM) that predicts each subsequent token conditioned on all preceding tokens. Latent reasoning introduces a set of unobserved intermediate steps,  $Z = \{z_i\}_{i=1}^M$ , which act as a substitute for the explicit reasoning trace C (see Fig. 2). The latent reasoning sequence begins with a special

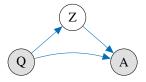


Figure 2: Graphical model of the latent reasoning generative model. The question prompt is used to generate continuous latent thought Z. The answer tokens are generated from the question and latent reasoning trace.

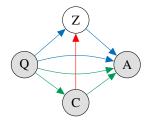


Figure 3: During training the student predicts the answer using latent tokens, teacher has the access to the full reasoning trace, and KV matching distills the information from the full to the latent CoT.

token <br/> <br/>bot>, continues with M continuous tokens, and terminates with <eot>, marking the end<br/> of the reasoning stage. During inference, these continuous latent tokens are generated by the same<br/> autoregressive model, bypassing the mapping of the embeddings to hard tokens. Instead, a (trainable) projection layer maps these continuous embeddings to the input embeddings that are used to predict the next token. We use the terms latent CoT and Continuous CoT (CCoT) interchangeably throughout the paper to refer to the tokens from Z.

**Training Objective.** Unlike chain-of-thought (CoT) reasoning traces, latent reasoning lacks direct supervision because latent traces are unobserved during training. Consequently, its performance is typically inferior to models trained with full CoT supervision (Deng et al., 2023; 2024). To address this, we leverage the observed reasoning traces C to guide latent reasoning during training, as illustrated in Fig. 3. This guidance is realized through distillation from teacher to student. Following Shen et al. (2025), we adopt a self-supervised framework in which the same model learns from explicit reasoning traces (as the teacher) as well as latent tokens (as the student).

We introduce KAVA, model with a novel objective, *KV-cache distillation*, to transfer relevant information from the teacher's reasoning trace to the student. An overview of this approach is depicted in Figure 1, with details provided in Section 3.2.

Our proposed KV-cache distillation loss is complementary to the CODI distillation loss introduced by Shen et al. (2025). CODI uses a single distillation token and matches its hidden activations between the teacher and the student models:

$$\mathcal{L}_{\text{CODI}} = \frac{1}{L} \sum_{l=1}^{L} \left| \text{sg}[\mathbf{h}_{t}^{l}] - \mathbf{h}_{s}^{l} \right|, \tag{1}$$

where L is the total number of layers in the model, sg is a stop-gradient operator and  $h^l$  are model's hidden activation from layer l. The distillation token is chosen as the one preceding the answer. For example, if the answer is formatted as "The answer is:5", the semicolon ":" is used as the distillation token.

We combine KV-cache distillation with the CODI self-distillation to add a richer supervision signal to the latent reasoning trace. The total training objective is the following:

$$\mathcal{L}_{\text{KaVa}} = -\underbrace{\frac{1}{N_A} \log p(A|Z,Q)}_{\text{Student loss}} - \underbrace{\frac{1}{N_A + N_C} \log p(A,C|Q)}_{\text{Teacher loss}} + \underbrace{\alpha_1 \mathcal{L}_{\text{CODI}} + \alpha_2 \mathcal{L}_{\text{KV}}}_{\text{CODI and KV distillation}}, \tag{2}$$

where  $\log p(\cdot)$  stands for cross-entropy loss,  $\alpha_1$  and  $\alpha_2$  are the hyperparameters that are used to balance the distillation terms,  $N_A$  and  $N_C$  denote number of tokens in the answer and CoT trace.

**Parallel Decoding.** Since latent tokens are generated sequentially, they do not allow for parallel decoding during training, which limits scalability. To mitigate this issue, we use Jacobi iteration over latent tokens to improve training and inference efficiency as proposed by Wu et al. (2025). Instead of generating latent tokens one by one during training PCCoT performs iterative updates of all tokens simultaneously for a predefined number of iterations T. PCCoT uses T < M, so that total number of forward passes is reduced from the number of latent tokens M to the number of iterations T. For T = M the method recovers the CODI explicitly and for T = 0 it corresponds to the Pause Token (Goyal et al., 2024).

#### 3.2 KV-CACHE DISTILLATION

To provide an additional supervision signal from the full chain-of-thought (CoT) trace to the latent reasoning process, KAVA uses a distillation method based on matching the respective key-value (KV) caches (last term in Eq. 2). We apply redundancy-aware KV-cache compression to the teacher's cache prior to distillation. This encourages the student to generate compressed and abstract representations, while preserving crucial reasoning information from the CoT trace.

We first extract the KV-cache for both the explicit reasoning trace (teacher) and the latent thought (student). Each cache consists of key and value tensors for every token i, layer  $l \in (1, \ldots, L)$ , and attention head  $h \in (1, \ldots, H)$  of the transformer:

$$K_t, V_t \in \mathbb{R}^{N_C \times H \times L \times d},$$
  

$$K_s, V_s \in \mathbb{R}^{M \times H \times L \times d},$$
(3)

where t stands for teacher and s for the student. We use the last Jacobi iteration T to extract the KV-cache of the student.

Addressing the Length Mismatch. The teacher cache  $(K_t,V_t)$  and student cache  $(K_s,V_s)$  differ in sequence length, since  $M < N_C$ . To align them while enforcing compression, we apply redundancy-aware KV eviction (Park et al., 2025; Cai et al., 2025) and obtain a compressed teacher cache  $\widetilde{K}_t, \widetilde{V}_t \in \mathbb{R}^{M \times H \times L \times d}$ . Specifically, we adapt R-KV (Cai et al., 2024) to select the top M KV-pairs (see App. A) based on a combined redundancy–importance score  $S_{i,h,l}$ :

$$S_{i,h,l} = \lambda \underbrace{I_{i,h,l}}_{\text{Importance}} + (1 - \lambda) \underbrace{R_{i,h,l}}_{\text{Redundancy}}, \quad \lambda \in [0, 1], \tag{4}$$

where  $\lambda$  is a hyperparameter controlling the balance between redundancy and importance. The eviction method is only applied during training, since the student is distilled to generate the compressed KV-cache. Since eviction method is not applied during inference, we leverage the answer tokens from the training data for the importance score computation. For each layer and head, we compute the attention score using the teacher's keys  $K_t^{\cdot,h,l} \in \mathbb{R}^{N_C \times d}$  and queries corresponding to the answer tokens tokens  $Q^{\cdot,h,l} \in \mathbb{R}^{N_A \times d}$ :

$$A^{\cdot,\cdot,h,l} = \operatorname{softmax}(Q^{\cdot,h,l} \cdot (K_t^{\cdot,h,l})^T) / \sqrt{d}) \in \mathbb{R}^{N_A \times N_C}. \tag{5}$$

The importance score is then aggregated over all answer tokens<sup>1</sup>:

$$I_{\cdot,h,l} = \frac{1}{N_A} \sum_{j} A^{j,\cdot,h,l} \in \mathbb{R}^{N_C}. \tag{6}$$

Note that this computation incurs negligible overhead, since the attention scores were computed during the teacher's forward pass. Following R-KV<sup>2</sup>, we compute a redundancy score  $R_{i,h,l}$  as the average pairwise cosine similarity among all key vectors and normalize via softmax.

Finally, we use the score values  $S_{i,h,l}$  (Eq. 4) to select top-M keys (and their corresponding values) for each head and layer in the teacher's KV-cache. Full details and pseudocode are provided in App. A.

**KV Matching.** Independent KV-pair eviction across layers and heads alters the cache's structure and contents, yet it remains usable by the original model (see Figure 1b). However, there no longer exists a correspondence between the resulting cache and hard tokens. For that reason, we cannot apply standard ways of distillation, matching the activations of the teacher and student model. Instead, we propose distilling the keys and values directly.

To this end, we distill the latent reasoning cache to match the compressed teacher's cache, effectively guiding the latent model to approximate the full reasoning process in a more efficient and abstract form. We combine the loss for the keys and values in equal weights to get the final term of Eq. 2:

$$\mathcal{L}_{KV} = \frac{1}{2M} \left( \|\mathbf{sg}[\widetilde{K}_t] - K_s\|_p^p + \|\mathbf{sg}[\widetilde{V}_t] - V_s\|_p^p \right), \tag{7}$$

<sup>&</sup>lt;sup>1</sup>For the group-query attention setting multiple queries are sharing the same key-value pair. In this case we apply MaxPool operation over the group before computing the importance score.

<sup>&</sup>lt;sup>2</sup>Official R-KV implementation is available at https://github.com/Zefan-Cai/R-KV.

Table 1: Test accuracy on in-distribution test dataset and zero-shot evaluation on out-of-distribution datasets. We use  $\dagger$  to denote results copied from Shen et al. (2025) and Wu et al. (2025). We consider full CoT as an upper bound on the performance and denote **best** latent reasoning method in bold and second-best with the line. We denote out method as KAVA.

Method		GSM8k-AUG		GSM8k-AUG-NL						
	GSM8k	GSM8k-Hard	SVAMP	GSM8k	GSM8k-Hard	SVAMP				
		Qw	VEN2.5 - 0.5	B - INSTRUCT						
FULL COT	50.6	12.6	54.3	48.5	12.6	57.3				
No-CoT	31.5	7.4	34.5	<u>31.5</u>	7.4	34.5				
CODI	<u>37.5</u>	8.1	<u>47</u>	20.2	4.9	33.3				
PCCoT	20.5	4.1	33	19.1	4.2	30.2				
KAVA (ours)	46.9 (1.4)	10.8 (0.1)	50.6 (0.4)	44.4 (1.8)	10.2 (0.4)	46.5 (0.1)				
		LL	AMA3.2 - 1	B - INSTRUCT						
FULL COT	61.6 <sup>†</sup>	15.6 <sup>†</sup>	66.7 <sup>†</sup>	53.2	13.3	62.9				
No-CoT	30.9 <sup>†</sup>	7.1 <sup>†</sup>	44.1 <sup>†</sup>	33.1	7.7	41.4				
<b>ICOT</b>	$19.0^{\dagger}$	$4.4^{\dagger}$	40.9 <sup>†</sup>	15.2 <sup>†</sup>	-	-				
COCONUT	$45.3^{\dagger}$	9.9 <sup>†</sup>	48.8 <sup>†</sup>	27.2 <sup>†</sup>	-	-				
CODI	55.6 <sup>†</sup>	$12.8^{\dagger}$	61.1 <sup>†</sup>	49.7 <sup>†</sup>	-	-				
PCCoT	53.35 <sup>†</sup> (0.18)	-	-	50.72 <sup>†</sup> (1.39)	-	-				
KAVA (ours)	56.5 (0.4)	<u>12.7</u> (0.1)	<u>58.9</u> (0.5)	55.7 (0.4)	12.8 (0.2)	58.6 (0.3)				
		LL	AMA3.2 - 3	B - INSTRUCT						
FULL COT	73.2	21.6	78.0	68.4	20.5	77.6				
No-CoT	41.7	10.5	56.9	41.7	10.5	56.9				
CODI	<u>61.0</u>	<u>15.0</u>	72.4	<u>55.9</u>	<u>13.6</u>	70.1				
PCCoT	54.7	13.5	69.5	47.6	11.0	65.2				
KAVA (ours)	65.7	15.2	<u>72.7</u>	60.0	14.8	<u>66.1</u>				

where  $\|\cdot\|_p$  denotes an  $L^p$ -norm. That is, we have  $L_1$  loss for p=1 and MSE loss for p=2. Note, that we first generate the whole student sequence with Jacobi iterations and then perform the distillation.

#### 4 EXPERIMENTS

#### 4.1 SETUP

We follow the experimental setup of Shen et al. (2025) and Wu et al. (2025) and extend the evaluation to more LLM families. Below we discuss the setup in more detail.

Model. We conduct experiments using the pretrained LLaMA3.2-1b-Instruct, LLaMA3.2-3b-Instruct and Qwen2.5-0.5b-Instruct (Grattafiori et al., 2024; Team, 2024) models and fine-tune them using LoRA (Hu et al., 2022). We follow Shen et al. (2025) and Wu et al. (2025) by using the same LoRA setup (rank 128 with alpha value 32 and dropout 0.1) for all the experiments. We employ PCCoT, the approach proposed by Wu et al. (2025), to generate latent thoughts; where 24 continuous latent tokens are generated in parallel with 3 iterations.

We fine-tune the models on two datasets: GSM8k-AUG, GSM8k-AUG-NL (Deng et al., 2023). Both datasets are augmented versions GSM8k (Cobbe et al., 2021), containing 385k training examples, with traces generated by GPT-4. GSM8k-AUG is then further processed by keeping only *equations* and removing all natural language from the traces. We provide a detailed description of the datasets in Appendix B. For in-distribution evaluation, we assess all models on the test split of the original GSM8k dataset (Cobbe et al., 2021). For zero-shot evaluation, we assess model generalization on two benchmarks: GSM8k-Hard (Gao et al., 2023) and SVAMP (Patel et al., 2021).

**Hyperparameters.** For our method, we conduct a hyperameter sweep over the learning rate, KV-cache distillation loss coefficient  $(\alpha_2)$ ,  $L^p$  norm of the loss and the normalization method (layer-wise loss normalization or none). We choose the best-performing model on validation and run this setting with three random seeds. We report all hyperparameters in Appendix C.

Table 2: We measure the efficiency of different reasoning model by the average number of forward passes required to generate the reasoning trace and answer. We use † to denote results copied from Shen et al. (2025) and Wu et al. (2025). We report the improvement in efficiency compared to the Full CoT in (parentheses).

Method		GSM8k-AUG		GSM8k-AUG-NL						
	GSM8k	GSM8k-Hard	SVAMP	GSM8k	GSM8k-Hard	SVAMP				
		Q'	WEN2.5 - 0.5	B - INSTRUC	Γ					
FULL COT	40.4	59.6	23.3	82.4	105.2	44.9				
No-CoT/ ICoT	7.4	10.1	7.0	7.4	10.1	7.0				
CODI	14.4	20.7	14.1	14.0	19.0	13.4				
KAVA (ours)	<u>9.5</u> (-76%)	<u>13.3</u> (-78%)	<u>8.9</u> (-62%)	<u>9.2</u> (-89%)	<u>13.5</u> (-87%)	<u>9.0</u> (-80%)				
		L	LAMA3.2 - 1	B - INSTRUCT						
FULL COT	65 <sup>†</sup>	-	-	71.9	80.2	40.6				
No-CoT / ICoT	-	-	-	6.2	7.3	6.2				
CODI	9†	-	-	-	-	-				
COCONUT	9†	-	-	_	-	-				
KAVA (ours)	<b>6.9</b> (-89%)	9.1	6.5	<u>7</u> (-90%)	<u>10</u> (-88%)	<u>6.4</u> (-86%)				
		L	LAMA3.2 - 31	BB - INSTRUCT						
FULL COT	31.6	40.3	17.0	75.2	32.9	38.3				
No-CoT / ICoT	6.1	7.4	<u>6.1</u>	6.1	7.4	<u>6.1</u>				
CODI	11.5	14.2	11.0	11.1	13.1	10.7				
KAVA (ours)	<u>6.4</u> (-80%)	<u>8.2</u> (-80%)	6 (-65%)	6 (-92%)	<u>7.9</u> (-76%)	<b>5.7</b> (-85%)				

We report the results of baseline approaches from Shen et al. (2025) and Wu et al. (2025) where possible. For the models not used in prior work, we take the hyperparameters from LLaMA3.2-1b, sweep over learning rates and report the result for the best performing model. We compare our method to CODI (Shen et al., 2025), PCCoT (Wu et al., 2025), Implicit CoT (iCoT) (Deng et al., 2024) and Coconut (Hao et al., 2024). We report the Full CoT performance as an upper bound and No-CoT as a lower bound.

#### 4.2 RESULTS

We report the average performance with standard error in Table 1. KAVA consistently outperforms the baselines. Importantly, we observe that KAVA has a lower drop in performance when switching from artificial GSM8k-AUG to a more realistic GSM8k-AUG-NL dataset. In the latter scenario, compression of the Full CoT trace would be more substantial as the traces are considerably longer, while questions are kept the same. This demonstrates the better scalability of out approach.

We also measure the efficiency of the method by the number of forward passes a model makes to generate the reasoning trace and the answer, reported in Table 2. KAVA builds on top of PCCoT, where we only use T=3 iterations (forward passes) to generate all the latent tokens. For that reason, we skip the PCCoT results in the table as they would be similar to ours. Our method achieves better efficiency than CoT, requiring between 62% and 92% fewer forward passes per question compared to Full CoT.

#### 4.3 ABLATION STUDIES

We select LLAMA3.2-1B-INSTRUCT to conduct ablation studies for our method. We run each experiment with three random seeds and report average test accuracy.

**Model Components.** First, we study how different modeling choices influence the final performance. In Table 3 we report benchmark performance when trained without the distillation loss (Shen et al., 2025) or without projection layer. As can be seen, both components are quite crucial, but even without them the method considerably outperforms the no-CoT baseline.

**Removing Last Step of the Trace.** Following Shen et al. (2025); Wu et al. (2025) we remove the last step from the teacher's reasoning trace. CODI demonstrates that this step is crucial for model performance, since otherwise the token that CODI chooses for distillation tends to be less informative. In Table 4 we train our model (using both KV matching and distillation) and PCCoT (only distillation) on all steps. Performance of our method drops much lower, indicating that KV-cache

Prj. ✓

X

 $\mathcal{L}_{KD}$ 

X

Table 3: Test accuracy on GSM8k dataset without projection layer and distillation loss ( $\alpha_1 = 0$ ).

GSM8k

56.5 (0.4)

52.8 (0.1)

52.2 (0.6)

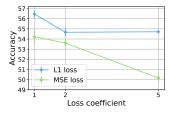
GSM-Hard	SVAMP
12.7 (0.1)	58.9 (0.5)

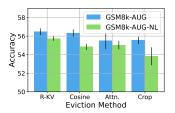
12.2(0.1)

12.3 (0.2)

Table 4: Test accuracy on GSM8k dataset when the teacher is trained on all the steps.

$\mathcal{L}_{KD}$	$\mathcal{L}_{KV}$	<b>Drop Last</b>	All Steps
1	1	56.5 (0.4)	51.2 (0.8)
1	X	53.35 (0.18)	47.2 (2.9)





56.2 (0.2)

58.3 (0.3)

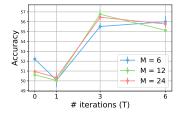


Figure 4: Test accuracy (%) of KAVA for different KV matching coefficient and loss.

Figure 5: Test accuracy (%) of KAVA with different eviction methods.

Figure 6: Test accuracy (%) of KAVA with different number of iterations and latent tokens.

distillation loss compensates for the lack of usefulness of a distillation token in a fully automatic manner.

**KV Loss Sensitivity.** Matching keys and values of the KV-cache is a non-standard way of distillation. Therefore, we study the model sensitivity to the distillation loss type and coefficient. In Figure 4 we plot the test accuracy for two losses and three different coefficients. The model performs consistently better with  $L_1$  loss when trained on GSM8k-AUG and with Llama-1b. However, we observed that better performance may be achieved when using MSE loss on other datasets (see Appendix C for the detailed hyperparameters used for all models and datasets).

**KV Eviction.** We follow Cai et al. (2025) in using  $\lambda=0.1$  (see Eq. 4) in R-KV eviction for all the experiments. As an ablation study we consider the two extremes: cosine-only ( $\lambda=0$ ) and attention-only ( $\lambda=1$ ). These cases correspond to choosing the keys and values based on diversity or importance only. Furthermore, we use a simple baseline of cropping the full CoT trace from the right, that is we only keep first M tokens of the teacher's cache for distillation. We report the results in Figure 5. We observe that combining both attention-based and similarity-based criteria enhances the performance for both datasets.

**Number of Tokens and Iterations.** Similarly to Wu et al. (2025), we observe that the number of iterations can have a different impact on accuracy depending on the number of latent tokens (Fig. 6). For larger numbers of latents (12, 24) we observe reduced performance beyond a certain number of iterations.

#### 5 INTERPRETABILITY OF LATENT REASONING TRACES

#### 5.1 DECODING THE LATENT TRACE

Although the latent CoT is not directly interpretable, one can still attempt to decode the reasoning trace from latent tokens. A straightforward approach is to project the final hidden state of the latent tokens via the language modeling head. An example of a decoded trace is shown in Table 5. More examples of the decoded traces are given in the Appendix E. Interestingly, the decoded latent trace is often identical to the trace generated by the teacher model, underlining the importance of the teacher guidance. In particular cases, as shown in the table, a reasoning step can be expressed in two equivalent forms (e.g. <<650\*2=1300>> and <<2\*650=1300>>). In regular CoT, this ambiguity is resolved after sampling a unique prefix of one of the variants, however, there is no explicit mechanism allowing for such resolution in a latent CoT. Nevertheless, the student arrives at the correct answer.

Table 5: Decoding the latent thoughts. A validation prompt is used: "Mrs. Taylor bought two smart televisions that cost \$650 each. If the total sales price had a 25% discount, how much did Mrs. Taylor pay for the two televisions?". Latent thoughts 16-24 are not shown due to their limited semantic value. 3 tokens with the highest logits are shown for each latent thought. Tokens T1, T2, T3, T4, T5, T6, T7 stand for \_total, \_cost, \_dollars, \_discount, \_original, \_gross, and \_price respectively. Following CODI, the teacher is trained on traces omitting the last step.

TopK	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	Answer
	GSM8K-Aug															
1 2 3	650 2 65	* + -	2 650 0	= * =\$	130 650 125	0 >> 00	>> I	<< The << (	_of _and	0 * k	* % *.	* % =	>> = 0	= * _=	= 325 125	975
Teacher	Teacher   <<650*2=1300>><<1300*25/100=325>>													975		
Golden	1		<<	650*2=	1300>>	<<130	0*25/1	100=325	>><<130	0-325=	=975>>	>				975
						GS	SM8K-A	Aug-NL								
1 2 3	T1 T2 T5	_of T2 T7	_of T2 _was	0 T3 T6	_\$ \$ _was	_\$ \$	_\$ \$	_\$ \$	_\$ \$	_\$ \$	_\$ \$	]\$ \$	T4 The	T4	T4	975
Teacher	Teacher   The total cost of the two televisions is 2 x \$650 = \$1300 [] \$1300 x 25/100 = \$325.													975		
Golden	The	total o	cost of	the tw	vo smart	tele	visio	ns is	[] \$	975 f	or the	e two	smart	televi	sions.	975

Models trained on the GSM8k-AUG dataset tend to produce latent CoT's that are easily interpretable. In contrast, models trained on the GSM8k-AUG-NL dataset resist this straightforward read-out method. We hypothesize that this is caused by the KV-cache distillation employed by KAVA —in a dataset with shorter traces, such as GSM8k-AUG, most of the time the KV-cache retains all of its content after eviction. On longer traces, such as the ones found in GSM8k-AUG-NL, not all content of the KV-cache is preserved, and, furthermore, each latent thought's distillation target may consist of keys and values originating from different tokens of the teacher's CoT. This can prevent latent thought to hard token correspondence from arising.

#### 5.2 TEACHER-STUDENT KV-CACHE CORRESPONDENCE

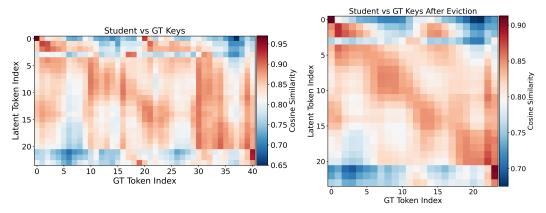


Figure 7: Cosine similarity of Keys in the latent CoT with Keys of the ground truth averaged across heads and layers. We use the same prompt and ground truth CoT as in Table 5.

We compute the cosine similarity of the keys and values in the latent CoT with (1) the ground truth KV-cache, and (2) the ground truth KV-cache after eviction. The results for Keys, averaged over attention heads and layers are presented in the Fig. 7. We observe that when comparing to the KV-cache after eviction, the similarities near the diagonal (x=y) tend to be higher, which is expected, as it is encouraged by the KV distillation. Furthermore, the values to the right of the diagonal are higher when comparing with the full CoT, which is desired, as this represents the compression of the original CoT (i.e. the key of some n-th latent token is similar to the key of an m-th hard token where n < m). The figure for Values and the full visualization of the similarities across layers and heads can be found in the Appendix D.

#### 6 CONCLUSION AND DISCUSSION

We introduce KAVA, a novel framework that bridges the supervision gap in latent reasoning by distilling knowledge from a teacher model's compressed Key-Value (KV) cache. Our central contribution is the demonstration that a compressed KV-cache, despite losing direct token correspondence, can serve as a rich, stepwise supervisory signal for a latent reasoning student. By aligning the student's latent trajectory with the teacher's internal reasoning dynamics in KV space, KAVA overcomes the limitations of token-level distillation and the inefficiencies of verbose Chain-of-Thought (CoT) traces. KAVA consistently outperforms strong latent reasoning baselines, scales effectively to larger backbones, and shows robust performance on natural-language reasoning datasets where prior methods often struggle. While the advancement of latent reasoning is linked to the availability of large-scale training data to instill novel reasoning dynamics, our work establishes compressed KV-cache distillation as a scalable and effective supervision technique for developing efficient and powerful reasoning models.

#### REFERENCES

Zefan Cai, Yichi Zhang, Bofei Gao, Yuliang Liu, Yucheng Li, Tianyu Liu, Keming Lu, Wayne Xiong, Yue Dong, Junjie Hu, et al. Pyramidkv: Dynamic kv cache compression based on pyramidal information funneling. *arXiv* preprint arXiv:2406.02069, 2024.

Zefan Cai, Wen Xiao, Hanshi Sun, Cheng Luo, Yikai Zhang, Ke Wan, Yucheng Li, Yeyang Zhou, Li-Wen Chang, Jiuxiang Gu, et al. R-kv: Redundancy-aware kv cache compression for training-free reasoning models acceleration. *arXiv* preprint arXiv:2505.24133, 2025.

Vivek Chari, Guanghui Qin, and Benjamin Van Durme. Kv-distill: Nearly lossless learnable context compression for llms. *arXiv preprint arXiv:2503.10337*, 2025.

Xinghao Chen, Anhao Zhao, Heming Xia, Xuan Lu, Hanlin Wang, Yanjun Chen, Wei Zhang, Jian Wang, Wenjie Li, and Xiaoyu Shen. Reasoning beyond language: A comprehensive survey on latent chain-of-thought reasoning. *arXiv preprint arXiv:2505.16782*, 2025.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Oihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, Bing Xue, Bingxuan Wang, Bochao Wu, Bei Feng, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Qu, Hui Li, Jianzhong Guo, Jiashi Li, Jiawei Wang, Jingchang Chen, Jingyang Yuan, Junjie Oiu, Junlong Li, J. L. Cai, Jiaqi Ni, Jian Liang, Jin Chen, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Liang Zhao, Litong Wang, Liyue Zhang, Lei Xu, Leyi Xia, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Meng Li, Miaojun Wang, Mingming Li, Ning Tian, Panpan Huang, Peng Zhang, Qiancheng Wang, Qinyu Chen, Qiushi Du, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, R. J. Chen, R. L. Jin, Ruyi Chen, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shengfeng Ye, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shuting Pan, S. S. Li, Shuang Zhou, Shaoqing Wu, Shengfeng Ye, Tao Yun, Tian Pei, Tianyu Sun, T. Wang, Wangding Zeng, Wanjia Zhao, Wen Liu, Wenfeng Liang, Wenjun Gao, Wenqin Yu, Wentao Zhang, W. L. Xiao, Wei An, Xiaodong Liu, Xiaohan Wang, Xiaokang Chen, Xiaotao Nie, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xinyu Yang, Xinyuan Li, Xuecheng Su, Xuheng Lin, X. Q. Li, Xiangyue Jin, Xiaojin Shen, Xiaosha Chen, Xiaowen Sun, Xiaoxiang Wang, Xinnan Song, Xinyi Zhou, Xianzu Wang, Xinxia Shan, Y. K. Li, Y. Q. Wang, Y. X. Wei, Yang Zhang, Yanhong Xu, Yao Li, Yao Zhao, Yaofeng Sun, Yaohui Wang, Yi Yu, Yichao Zhang, Yifan Shi, Yiliang Xiong, Ying He, Yishi Piao, Yisong Wang, Yixuan Tan, Yiyang Ma, Yiyuan Liu, Yongqiang Guo, Yuan Ou, Yuduan Wang, Yue Gong, Yuheng Zou, Yujia He, Yunfan Xiong, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuyang Zhou, Y. X. Zhu, Yanhong Xu, Yanping Huang, Yaohui Li, Yi Zheng, Yuchen Zhu, Yunxian Ma, Ying

- Tang, Yukun Zha, Yuting Yan, Z. Z. Ren, Zehui Ren, Zhangli Sha, Zhe Fu, Zhean Xu, Zhenda Xie, Zhengyan Zhang, Zhewen Hao, Zhicheng Ma, Zhigang Yan, Zhiyu Wu, Zihui Gu, Zijia Zhu, Zijun Liu, Zilin Li, Ziwei Xie, Ziyang Song, Zizheng Pan, Zhen Huang, Zhipeng Xu, Zhongyu Zhang, and Zhen Zhang. Deepseek-rl: Incentivizing reasoning capability in Ilms via reinforcement learning, 2025. URL https://arxiv.org/abs/2501.12948.
- Yuntian Deng, Kiran Prasad, Roland Fernandez, Paul Smolensky, Vishrav Chaudhary, and Stuart Shieber. Implicit chain of thought reasoning via knowledge distillation. *arXiv* preprint arXiv:2311.01460, 2023.
- Yuntian Deng, Yejin Choi, and Stuart Shieber. From explicit cot to implicit cot: Learning to internalize cot step by step. *arXiv preprint arXiv:2405.14838*, 2024.
- Harry Dong, Xinyu Yang, Zhenyu Zhang, Zhangyang Wang, Yuejie Chi, and Beidi Chen. Get more with less: Synthesizing recurrence with kv cache compression for efficient llm inference. In *ICML*, 2024. URL https://openreview.net/forum?id=uhHDhVKFMW.
- Yu Fu, Zefan Cai, Abedelkadir Asi, Wayne Xiong, Yue Dong, and Wen Xiao. Not all heads matter: A head-level KV cache compression method with integrated retrieval and reasoning. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=FJFVmeXusW.
- Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. Pal: Program-aided language models. In *International Conference on Machine Learning*, pp. 10764–10799. PMLR, 2023.
- Sachin Goyal, Ziwei Ji, Ankit Singh Rawat, Aditya Krishna Menon, Sanjiv Kumar, and Vaishnavh Nagarajan. Think before you speak: Training language models with pause tokens. In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=ph04CRkPdC.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. arXiv preprint arXiv:2407.21783, 2024.
- Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. Training large language models to reason in a continuous latent space, 2024. URL https://arxiv.org/abs/2412.06769.
- Edward J Hu, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Binyuan Hui, Jian Yang, Zeyu Cui, Jiaxi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Keming Lu, Kai Dang, Yang Fan, Yichang Zhang, An Yang, Rui Men, Fei Huang, Bo Zheng, Yibo Miao, Shanghaoran Quan, Yunlong Feng, Xingzhang Ren, Xuancheng Ren, Jingren Zhou, and Junyang Lin. Qwen2.5-coder technical report, 2024. URL https://arxiv.org/abs/2409.12186.
- Junyoung Park, Dalton Jones, Matthew J Morse, Raghavv Goel, Mingu Lee, and Chris Lott. Keydiff: Key similarity-based kv cache eviction for long-context llm inference in resource-constrained environments. *arXiv preprint arXiv:2504.15364*, 2025.
- Arkil Patel, Satwik Bhattamishra, and Navin Goyal. Are nlp models really able to solve simple math word problems? In Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pp. 2080–2094, 2021.
- Jacob Pfau, William Merrill, and Samuel R. Bowman. Let's think dot by dot: Hidden computation in transformer language models. In First Conference on Language Modeling, 2024. URL https: //openreview.net/forum?id=NikbrdtYvG.

- Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, Hugh Zhang, Chen Bo Calvin Zhang, Mohamed Shaaban, John Ling, Sean Shi, Michael Choi, Anish Agrawal, Arnav Chopra, Adam Khoja, Ryan Kim, Richard Ren, Jason Hausenloy, Oliver Zhang, Mantas Mazeika, Dmitry Dodonov, Tung Nguyen, Jaeho Lee, Daron Anderson, Mikhail Doroshenko, Alun Cennyth Stokes, Mobeen Mahmood, Oleksandr Pokutnyi, Oleg Iskra, Jessica P. Wang, John-Clark Levin, Mstyslav Kazakov, Fiona Feng, Steven Y. Feng, Haoran Zhao, Michael Yu, Varun Gangal, Chelsea Zou, Zihan Wang, Serguei Popov, Robert Gerbicz, Geoff Galgon, Johannes Schmitt, Will Yeadon, Yongki Lee, Scott Sauers, Alvaro Sanchez, Fabian Giska, Marc Roth, Søren Riis, Saiteja Utpala, Noah Burns, Gashaw M. Goshu, Mohinder Maheshbhai Naiya, Chidozie Agu, Zachary Giboney, Antrell Cheatom, Francesco Fournier-Facio, Sarah-Jane Crowson, Lennart Finke, Zerui Cheng, Jennifer Zampese, Ryan G. Hoerr, Mark Nandor, Hyunwoo Park, Tim Gehrunger, Jiaqi Cai, Ben McCarty, Alexis C Garretson, Edwin Taylor, Damien Sileo, Qiuyu Ren, Usman Qazi, Lianghui Li, Jungbae Nam, and John B. Wydallis et al. Humanity's last exam, 2025. URL https://arxiv.org/abs/2501.14249.
- Utkarsh Saxena, Gobinda Saha, Sakshi Choudhary, and Kaushik Roy. Eigen attention: Attention in low-rank space for KV cache compression. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (eds.), *Findings of the Association for Computational Linguistics: EMNLP 2024*, pp. 15332–15344, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-emnlp.899. URL https://aclanthology.org/2024.findings-emnlp.899/.
- Zhenyi Shen, Hanqi Yan, Linhai Zhang, Zhanghao Hu, Yali Du, and Yulan He. Codi: Compressing chain-of-thought into continuous space via self-distillation, 2025. URL https://arxiv.org/abs/2502.21074.
- DiJia Su, Hanlin Zhu, Yingchen Xu, Jiantao Jiao, Yuandong Tian, and Qinqing Zheng. Token assorted: Mixing latent and text tokens for improved language model reasoning, 2025. URL https://arxiv.org/abs/2502.03275.
- Qwen Team. Qwen2 technical report. arXiv preprint arXiv:2407.10671, 2, 2024.
- Haoyi Wu, Zhihao Teng, and Kewei Tu. Parallel continuous chain-of-thought with jacobi iteration. *arXiv preprint arXiv:2506.18582*, 2025.
- Yige Xu, Xu Guo, Zhiwei Zeng, and Chunyan Miao. SoftCoT: Soft chain-of-thought for efficient reasoning with LLMs. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 23336–23351, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-251-0. doi: 10.18653/v1/2025.acl-long.1137. URL https://aclanthology.org/2025.acl-long.1137/.
- Ziyin Zhang, Jiahao Xu, Zhiwei He, Tian Liang, Qiuzhi Liu, Yansi Li, Linfeng Song, Zhenwen Liang, Zhuosheng Zhang, Rui Wang, Zhaopeng Tu, Haitao Mi, and Dong Yu. Deeptheorem: Advancing Ilm reasoning for theorem proving through natural language and reinforcement learning, 2025. URL https://arxiv.org/abs/2505.23754.
- Rui-Jie Zhu, Tianhao Peng, Tianhao Cheng, Xingwei Qu, Jinfa Huang, Dawei Zhu, Hao Wang, Kaiwen Xue, Xuanliang Zhang, Yong Shan, et al. A survey on latent reasoning. *arXiv preprint arXiv:2507.06203*, 2025.

#### A KV EVICTION DETAILS

We provide pseudocode to compute the r-KV score in Listing 1. The function takes as input a key-value pair and the attention scores between the CoT and and Answer tokens. There are several implementation differences from the original R-KV method.

**Padding Tokens** First, we need to take into account padding tokens since we evict KV-cache in a batch during training. We do that by always assigning the lowest possible redundancy and importance score to the value-key pairs corresponding to the padding tokens

**Importance Score** To compute the importance score, we use the attention score that answer tokens get when attending to the full CoT. We extract those value during the normal teacher forward pass and reuse to compute the

**Retention of Recent Tokens** R-KV implementation adjust the redundancy score by always keeping  $\beta$  the most recent tokens. This is important for a reliable model performance during generation. We only use our method during training and apply it to the whole reasoning trace, therefore we skip this adjustment and only rely on selecting the most diverse keys with high attention to the answer tokens.

Listing 1: Pseudocode to implement the eviction score for a given key-value pair.

```
def r_kv_score(key: torch.tensor, attn: torch.tensor, lbd: float):
1
2
3
       key: torch.tensor [bs, N_c, d] - CoT keys for a single head and layer
       attn: torch.tensor [bs, N_A, N_c] - attenton scores
4
       lbd: float - the weight of the importance score
5
6
       # compute redundancy score
       key_norm = key / (key.norm(dim=-1, keepdim=True) + 1e-8)
8
       cosine_sim = torch.einsum("...id,...jd->...ij", key_norm, key_norm)
9
       for i in range(cosine_sim.shape[0]):
10
           cosine_sim[i].fill_diagonal_(0)
11
12
       cos_score = torch.sum(-cosine_sim, dim=-2) / torch.sum(
           ~pad_tokens, dim=-1, keepdim=True
13
14
15
       # Normalize to 1
       R = cos_score.softmax(dim=-1)
16
17
       pad\_tokens = key.sum(-1) == 0
       R[pad\_tokens] = 0
18
19
       # compute importance score
20
21
       # sofmax over CoT dimention and avrage over answer tokens
       I = F.softmax(attn, dim=-1).mean(-2)
22
       # Assign the lowest score to the padding tokens
23
       I[pad\_tokens] = 0
24
25
26
       S = 1bd * I + (1 - 1bd) * R
27
28
       return S
```

#### B DATASETS

Our models are trained using the GSM8k-Aug and GSM8k-Aug-NL datasets introduced by Deng et al. (2023), which augment the training set of the GSM8k (Cobbe et al., 2021) using GPT4 and provide a separate validation split. The golden traces in the datasets are split into discrete steps. GSM8k-Aug traces consist only of succinct statements such as  $<<600 \times 30/100 = 180>>$ ;  $<<600 \times 10/100 = 60>>$ . The questions and answers in the NL (Natural Language) subset are identical, however the steps are formulated in natural language:  $600 \times 30/100 = 180$  employees were promoted.;  $600 \times 10/100 = 60$ 0 employees received a bonus.

	GSM8K-Aug	GSM8K-Aug-NL						
Huggingface Path	whynlp/gsm8k-aug	whynlp/gsm8k-aug-nl						
No. of Train Sample	385,620							
No. of Validation Samples	500							
No. of Test Samples	1	319						

# C HYPERPARAMETERS

Table 6: All the hyperparameters used for our method.

Hyperparameter		GSM8k-AUG-NL
	LLAMA3.2	- 1b - Instruct
$\overline{\alpha_1 \text{ (CODI)}}$	10	10
KV loss	Smooth L1	MSE
Layer-wise std	True	True
$\alpha_2$ (KV)	1	1
$\alpha_2(\mathbf{K}\mathbf{V})$ r-kv $\lambda$	0.1	0.1
Use Projection	True	True
learning rate	8e-4	8e-4
lr scheduler	Cosine	Cosine
	AdamW	AdamW
optimizer		
batch size	128	128
weight decay	0.1	0.1
gradient clipping	2	2
epochs	10	10
		0.5B - INSTRUCT
$\alpha_1$ (CODI)	10	10
KV loss	MSE	MSE
Layer-wise std	False	True
$\alpha_2$ (KV)	1	1
r-kv λ	0.1	0.1
Use Projection	True	True
learning rate	5e-4	8e-4
lr scheduler	Cosine	Cosine
optimizer	AdamW	AdamW
batch size	128	128
weight decay	0.01	0.1
gradient clipping	2	2
epochs	10	10
	LLAMA3.2	- 3b - Instruct
$\alpha_1$ (CODI)	20	20
KV loss	Smooth L1	Smooth L1
Layer-wise std	False	False
$\alpha_2$ (KV)	2	2
$r-kv \lambda$	0.1	0.0
Use Projection	True	False
learning rate	2e-4	2e-4
lr scheduler	Cosine	Cosine
optimizer	AdamW	AdamW
batch size	128	128
weight decay	0.1	0.1
gradient clipping	2	2
epochs	5	5
epociis	ر	<u> </u>

# D KV-CACHE COSINE SIMILARITY BETWEEN THE LATENT COT AND THE GROUND-TRUTH COT

We investigate the similarity between the KV-cache representing the latent CoT and the KV-cache of the ground-truth CoT. Figures 8 and 7 present the similarities averaged over layers and heads, while figures 9, 10, 11, and 12 show the similarities in individual heads and layers.

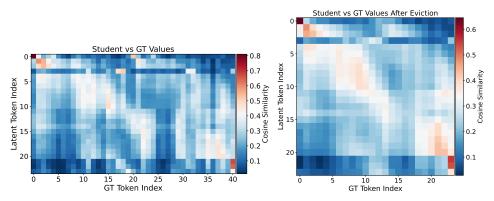


Figure 8: Cosine similarity of Values in the latent CoT with Values of the ground truth averaged across heads and layers. We use the same prompt and ground truth CoT as in Table 5.

Figure 9: Cosine similarity between Keys in the latent CoT and Keys of the ground truth across layers.

## E DECODED LATENT TRACES

In this section we present two additional examples of traces decoded in the same manner as described in section 5.1.

Student vs GT Values - Detailed

Figure 10: Cosine similarity between Values in the latent CoT and Values of the ground truth across layers.

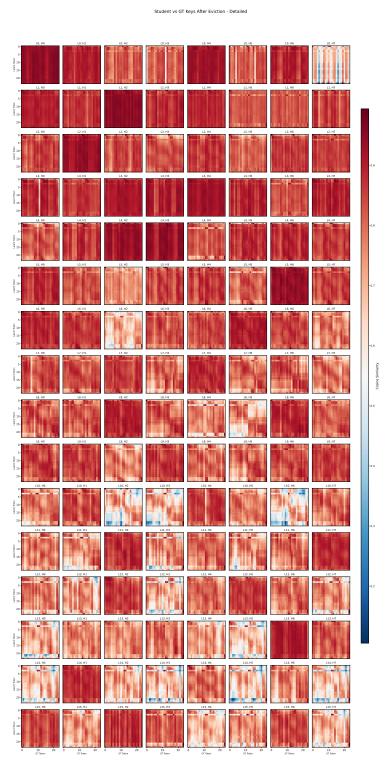


Figure 11: Cosine similarity between Keys in the latent CoT and Keys of the ground truth after eviction across layers.

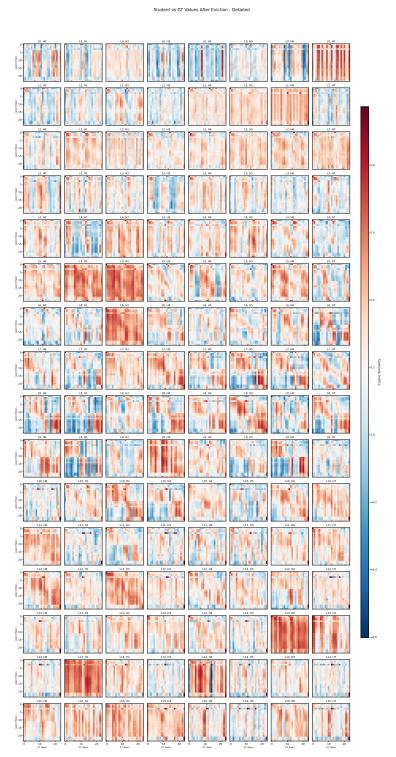


Figure 12: Cosine similarity between Values in the latent CoT and Values of the ground truth after eviction across layers.

TopK	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	Answer
	GSM8K-Aug															
1 2 3	24 50	* *. *(	50 0 30	= * *.	120 150 600	0 >> 00	>> <<	<< The << (	120 0	* *. 0	0 *•	0 10 00	0 >> 00	= >> 0	= 0 >>	3600
Teacher						<<5	0*0.10=	=5>><<5	*24=120	)>>						3600
Golden	1				<<50	*.10=5	>><<5 * 2	24=120>	><<120	×30=36	500>>					3600
							GSM	8K-Aug-l	NL							
1 2 3	T6 T7 T8	I6	50 0 *	T9 0 *	* T11	* T11	* T11	* T11	* ,	, *	, T11	, T10 _per	, T10 _per	, T10	0 _per 00	3600
Teacher	1					Не	gets 0	.10*50=	5 doll	ars a	hour					1800
Golden	Golden   He makes 50*\$.10=\$5 per hour [] \$120*30=\$3600 a month												3600			

Table 7: Prompt: "Jon runs a website where he gets paid for every person who visits. He gets paid \$0.10 for every person who visits. Each hour he gets 50 visits. His website operates 24 hours a day. How many dollars does he make in a 30 day month?". T6 – T11 stand for \_gets, \_makes, \_operates, \_visits, \_hourly, and \_hour respectively. Tokens 16-24 are omitted due to low semantic content.

TopK	1	2	3	4	5	6	7	8	9	10	11	12	13	14	Answer
GSM8K-Aug															
1 2 3	150 2 300	* + _*	2 1 5	= * _=	300 150 30	>> >>	The << T16	_as T15 _of	1500						
Teacher						<<1	50*2=30	00>>						1	1500
Golden					<<15	0*2=30	0>><<30	00*5=15	00>>						1500
						GSM	8K-Aug-l	NL							
1 2 3	T13 T11 T14	T11 _to T18	T11 T14 _to	T17 T12 T11	T11 _to T14	T11 T14 to	T11 T14	T11 T14	T11 T14	T11 T14	T11 T14	T11 T14	T11 T14	T11 T14	1500
Teacher	Teacher   Raine takes 150 x 2 = 300 steps walking to and from school in one day.												1500		
Golden	Rain	e takes	150 x	2 = 3	00 step	s walk	ing [.	] he	r 300	x 5 =	1500 st	eps in	five	days.	1500

Table 8: Prompt: "Raine's house is just a walking distance from her school. It takes her 150 steps to walk to the school. How many steps does she take walking to and from school in five days?". T11 - T18 stand for \_walking, \_footsteps, \_walks, \_walk, \_but, This, \_steps, and \_going respectively.